

SAN Assignment - regression

Martin Endler endlemar@fel.cvut.cz

Submission

Fill in your name above for clarity. To solve this homework, simply write your answers into this document and fill in the marked pieces of code. Submit your solution consisting of both this modified Rmd file and a knitted PDF document as an archive to the courseware BRUTE upload system for the SAN course. The deadline is specified there.

Initialization

Load the required libraries `gtools`, `caret` and `glmnet`, make sure you have those installed. We also fix the random seed for reproducibility convenience.

```
require(gtools);

## Loading required package: gtools
require(caret);

## Loading required package: caret
## Loading required package: ggplot2
## Loading required package: lattice
require(glmnet);

## Loading required package: glmnet
## Loading required package: Matrix
## Loaded glmnet 4.1-4
##
## Attaching package: 'glmnet'
## The following object is masked from 'package:gtools':
##
##   na.replace
set.seed(0)
```

Here, we define constants of the assignment. You may play with the values and observe what happens, but in your solution you should use the given values unchanged.

```
n.samples <- 256 # Total number of samples (training and testing together)
n.dimensions <- 100 # Number of n.dimensions, a.k.a. attributes or features
```

Model evaluation procedure

The function `learnAndTest` takes a matrix of independent variables values `X` and a corresponding vector of dependent variable (a.k.a. response) `Y` then trains and evaluates a model specified by the `modelType` parameter.

If you are interested in all possible `modelType` parameter values, refer to <http://topepo.github.io/caret/train-models-by-tag.html>. Here we will only be using three: "lm" for ordinary least squares and "glmnet" with parameter `alpha = 1` resp. `alpha = 0` for LASSO resp. Ridge. (There are also `lasso` and `ridge` methods you may try, but these have inconsistent API for passing `lambda`.) This function learns the model from the data and estimates its accuracy using cross validation. The results are then printed to output. For purpose of this assignment, consider only the RMSE (root-mean-square error) criterion.

```
learnAndTest <- function(X, Y, modelType, ...) {
  data <- data.frame(X, Y)

  train.control <- trainControl(method = "cv", number = 10) # 10-fold cross-validation
  # Alternative to get more accurate, but slower evaluation:
  # train.control <- trainControl(method = "LOOCV")

  # Here we train the model using the versatile `train` function from the caret package
  train(Y ~ .,
        data = data,
        method = modelType,
        trControl = train.control,
        ...
  )
}
```

We will also precompute an array of candidate `lambda` values for LASSO and Ridge.

```
lambda_lasso <- expand.grid(lambda = 10^seq(10, -3, length = 10), alpha = 1)
lambda_ridge <- expand.grid(lambda = 10^seq(10, -3, length = 10), alpha = 0)
```

Initial data generation

Here, we generate some data.

```
# Generates independent variables by uniform i.i.d. sampling
X <- replicate(
  n.dimensions,
  runif(n.samples, min = -10, max = 10)
)

# Randomly generates the actual underlying coefficients of the linear dependency
coefs <- runif(n.dimensions, min = 1, max = 4)
intercept <- 0 # For simplicity

# Synthesizes dependent variable (observed values) by the given linear dependency plus noise
noise <- rnorm(n.samples, sd = 8, mean = 0) # Gaussian noise to be added to the response
Y <- (X %*% coefs) + intercept + noise # Note: (%*%) is the matrix multiplication operator
```

Testing the models

Now let us run the following tests:

```
print(learnAndTest(X, Y, "lm"))
```

```
## Linear Regression
##
## 256 samples
## 100 predictors
```

```

##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 231, 231, 230, 231, 231, 229, ...
## Resampling results:
##
##      RMSE      Rsquared    MAE
##      10.22282  0.9954098  8.107877
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
print(learnAndTest(X, Y, "glmnet", tuneGrid = lambda_ridge))

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## glmnet
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 229, 231, 231, 231, 228, 230, ...
## Resampling results across tuning parameters:
##
##      lambda      RMSE      Rsquared    MAE
##      1.000000e-03  13.58962  0.9933283  11.09131
##      2.782559e-02  13.58962  0.9933283  11.09131
##      7.742637e-01  13.58962  0.9933283  11.09131
##      2.154435e+01  34.01922  0.9740668  27.24431
##      5.994843e+02  120.12106  0.7891572  97.88858
##      1.668101e+04  143.78427  0.7190257  117.35115
##      4.641589e+05  144.96674      NaN  118.32455
##      1.291550e+07  144.96674      NaN  118.32455
##      3.593814e+08  144.96674      NaN  118.32455
##      1.000000e+10  144.96674      NaN  118.32455
##
## Tuning parameter 'alpha' was held constant at a value of 0
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 0 and lambda = 0.7742637.
print(learnAndTest(X, Y, "glmnet", tuneGrid = lambda_lasso))

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## glmnet
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 230, 230, 230, 231, 228, 232, ...
## Resampling results across tuning parameters:

```

```
##
##      lambda      RMSE      Rsquared      MAE
##  1.000000e-03   10.43912  0.9950887    8.530893
##  2.782559e-02   10.43912  0.9950887    8.530893
##  7.742637e-01   21.50918  0.9851375   16.937733
##  2.154435e+01  135.67380  0.1978032  109.662368
##  5.994843e+02  145.51104      NaN    118.523560
##  1.668101e+04  145.51104      NaN    118.523560
##  4.641589e+05  145.51104      NaN    118.523560
##  1.291550e+07  145.51104      NaN    118.523560
##  3.593814e+08  145.51104      NaN    118.523560
##  1.000000e+10  145.51104      NaN    118.523560
##
## Tuning parameter 'alpha' was held constant at a value of 1
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 1 and lambda = 0.02782559.
```

Task 1:

Answer the following questions:

- What change in the learned model would you anticipate if we changed the **mean** parameter value to a different constant in the noise generation? You may answer this either by talking about the coefficients or by giving a geometrical interpretation. **Answer:** From a geometric point of view, the predicted y values would be shifted up (for a positive mean), resp. down (for a negative mean). This shift would correspond to a change in the estimated intercept parameter. The other estimated coefficients would remain the same.
- In our example we generated samples (i.e. the independent variables **X**) from uniform distribution. The least squares method, on the other hand, has something called the “normality assumption”. Have we violated that assumption? Justify. **Answer:** No, we haven’t. The “normality assumption” requires that the residuals are normally distributed, not the data (our generated samples).
- Which method gave the best results? Is it the most common one to do so if you re-run the test several times? Why do you think it performs better the best? **Answer:** The LASSO gave the best result (smallest RMSE) most often while re-running the test multiple times. However, the ordinary Least squares method was also very good, with a small RMSE very close to the one from LASSO. From the very small lambda parameter in the best LASSO, we can see that the LASSO is almost the same as the Least squares method (the shrinkage penalty term is close to zero).
- Check the selected values for the lambda parameter for ridge and LASSO. Are they low or high? How does it relate to the above answer? **Answer:** They are both very low. That means that both ridge and LASSO are almost the same as the Least squares method (the shrinkage penalty term is close to zero).

Least squares assumptions

The data generation model assumed by the ordinary least squares method can be mathematically written as follows:

$$Y = \mathbf{X}^T \boldsymbol{\beta} + \beta_0 + G, \quad G \sim \mathcal{N}(0, \sigma^2)$$

This formula implicitly expresses some of the assumptions about the data, required for the method to work reliably.

- The observed value Y is influenced by some Gaussian noise G .
- There truly exists an underlying linear dependency.
- The noise is homoscedastic (σ^2 is a constant).

Task 2:

First of all, make sure you understand how elements of this formula correspond to the code in the “data generation” section.

Your task is to violate each of these assumptions (one at a time) and **briefly** comment the changes in the learned model by statistically comparing it to model using the above data. (Coefficient summary below.) The catch here is that you are allowed to only modify the noise generation procedure to achieve that. Attempt to find a way of violating the assumptions to achieve a clear difference, but any solution that is technically correct will be awarded full points.

It is sufficient to look at the **summary** of the OLS model.

```
summary(learnAndTest(X, Y, "lm"))
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-16.3220	-3.7072	0.1499	4.0184	16.0036

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.43023	0.63316	-0.679	0.498
X1	3.32876	0.12014	27.708	<2e-16 ***
X2	1.61195	0.10468	15.399	<2e-16 ***
X3	2.42437	0.11129	21.784	<2e-16 ***
X4	2.90811	0.09546	30.466	<2e-16 ***
X5	1.15971	0.10603	10.937	<2e-16 ***
X6	3.82874	0.10620	36.051	<2e-16 ***
X7	2.37577	0.11010	21.578	<2e-16 ***
X8	1.94196	0.10157	19.119	<2e-16 ***
X9	1.26780	0.10836	11.700	<2e-16 ***
X10	2.43546	0.11277	21.597	<2e-16 ***
X11	2.79039	0.11020	25.322	<2e-16 ***
X12	1.84899	0.09814	18.840	<2e-16 ***
X13	3.31907	0.10670	31.105	<2e-16 ***
X14	3.73201	0.10760	34.684	<2e-16 ***
X15	2.85848	0.10641	26.863	<2e-16 ***
X16	2.18332	0.10949	19.941	<2e-16 ***
X17	2.98273	0.10435	28.583	<2e-16 ***
X18	3.95243	0.11494	34.388	<2e-16 ***
X19	2.72129	0.11400	23.870	<2e-16 ***
X20	2.26276	0.10559	21.430	<2e-16 ***
X21	1.45329	0.11139	13.047	<2e-16 ***
X22	2.06841	0.10820	19.117	<2e-16 ***
X23	2.65462	0.10546	25.172	<2e-16 ***
X24	2.43848	0.11204	21.764	<2e-16 ***
X25	3.39725	0.10541	32.230	<2e-16 ***
X26	3.63669	0.10717	33.933	<2e-16 ***
X27	4.08163	0.12419	32.867	<2e-16 ***
X28	1.39751	0.10846	12.885	<2e-16 ***
X29	1.21262	0.10735	11.296	<2e-16 ***
X30	3.36342	0.10379	32.405	<2e-16 ***

## X31	1.98639	0.11157	17.804	<2e-16	***
## X32	2.06385	0.10752	19.194	<2e-16	***
## X33	3.41435	0.10517	32.465	<2e-16	***
## X34	1.04075	0.10473	9.938	<2e-16	***
## X35	2.57185	0.12316	20.883	<2e-16	***
## X36	1.44023	0.10446	13.787	<2e-16	***
## X37	3.28405	0.11017	29.809	<2e-16	***
## X38	2.72006	0.10675	25.480	<2e-16	***
## X39	2.26690	0.10584	21.418	<2e-16	***
## X40	1.88773	0.10029	18.823	<2e-16	***
## X41	2.95467	0.10517	28.094	<2e-16	***
## X42	2.47410	0.12403	19.948	<2e-16	***
## X43	3.85890	0.10951	35.236	<2e-16	***
## X44	1.37370	0.10343	13.281	<2e-16	***
## X45	1.31642	0.11187	11.767	<2e-16	***
## X46	1.89909	0.11339	16.749	<2e-16	***
## X47	1.35180	0.11030	12.255	<2e-16	***
## X48	2.36427	0.10937	21.617	<2e-16	***
## X49	3.06998	0.11053	27.776	<2e-16	***
## X50	1.69993	0.10494	16.199	<2e-16	***
## X51	2.66625	0.10163	26.236	<2e-16	***
## X52	2.27966	0.10663	21.380	<2e-16	***
## X53	2.46486	0.10828	22.763	<2e-16	***
## X54	2.02121	0.10601	19.065	<2e-16	***
## X55	4.06324	0.10403	39.058	<2e-16	***
## X56	3.48045	0.11650	29.875	<2e-16	***
## X57	3.82223	0.11656	32.792	<2e-16	***
## X58	1.83125	0.10265	17.841	<2e-16	***
## X59	2.07426	0.11293	18.368	<2e-16	***
## X60	1.01282	0.10589	9.564	<2e-16	***
## X61	1.88231	0.10576	17.798	<2e-16	***
## X62	2.19831	0.11279	19.490	<2e-16	***
## X63	3.36844	0.11272	29.884	<2e-16	***
## X64	1.62348	0.10764	15.082	<2e-16	***
## X65	2.63364	0.10530	25.010	<2e-16	***
## X66	2.26839	0.10769	21.064	<2e-16	***
## X67	3.15058	0.10654	29.572	<2e-16	***
## X68	2.11758	0.10611	19.957	<2e-16	***
## X69	3.48378	0.11229	31.026	<2e-16	***
## X70	2.20811	0.10557	20.915	<2e-16	***
## X71	2.45733	0.11585	21.211	<2e-16	***
## X72	3.93237	0.10572	37.195	<2e-16	***
## X73	3.03948	0.11007	27.615	<2e-16	***
## X74	2.77817	0.10538	26.363	<2e-16	***
## X75	2.04650	0.11603	17.638	<2e-16	***
## X76	2.18996	0.11704	18.711	<2e-16	***
## X77	1.84367	0.10785	17.095	<2e-16	***
## X78	3.64003	0.10593	34.361	<2e-16	***
## X79	2.09904	0.11699	17.943	<2e-16	***
## X80	2.31115	0.10571	21.863	<2e-16	***
## X81	1.60042	0.10789	14.833	<2e-16	***
## X82	1.34889	0.11418	11.814	<2e-16	***
## X83	2.51265	0.11094	22.649	<2e-16	***
## X84	2.70763	0.10672	25.371	<2e-16	***

```
## X85      1.17989    0.10689   11.038   <2e-16 ***
## X86      2.10964    0.10729   19.663   <2e-16 ***
## X87      2.16377    0.10097   21.431   <2e-16 ***
## X88      3.07060    0.11259   27.272   <2e-16 ***
## X89      2.72039    0.10465   25.995   <2e-16 ***
## X90      1.51493    0.12133   12.486   <2e-16 ***
## X91      2.63591    0.10543   25.002   <2e-16 ***
## X92      1.78357    0.10841   16.453   <2e-16 ***
## X93      1.66301    0.10155   16.377   <2e-16 ***
## X94      1.56552    0.10155   15.416   <2e-16 ***
## X95      1.47395    0.10822   13.620   <2e-16 ***
## X96      2.82689    0.11436   24.719   <2e-16 ***
## X97      2.07620    0.11338   18.311   <2e-16 ***
## X98      2.81395    0.10439   26.957   <2e-16 ***
## X99      1.97136    0.10571   18.649   <2e-16 ***
## X100     1.65748    0.10818   15.321   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.818 on 155 degrees of freedom
## Multiple R-squared:  0.9983, Adjusted R-squared:  0.9971
## F-statistic: 888.8 on 100 and 155 DF, p-value: < 2.2e-16
```

```
noise <- runif(n.samples, min = -8, max = 8)

# KEEP THE CODE BELOW
Y <- (X %*% coefs) + intercept + noise
summary(learnAndTest(X, Y, "lm"))
```

Violate noise normality

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.6597  -2.4261   0.1438   2.3403  10.0544
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.22102    0.36360   0.608   0.544
## X1           3.35392    0.06899  48.614 <2e-16 ***
## X2           1.75900    0.06012  29.260 <2e-16 ***
## X3           2.29553    0.06391  35.917 <2e-16 ***
## X4           2.79220    0.05482  50.936 <2e-16 ***
## X5           1.18499    0.06089  19.460 <2e-16 ***
## X6           3.86302    0.06099  63.339 <2e-16 ***
## X7           2.18476    0.06323  34.553 <2e-16 ***
## X8           1.95167    0.05833  33.459 <2e-16 ***
## X9           1.21902    0.06223  19.590 <2e-16 ***
## X10          2.60791    0.06476  40.270 <2e-16 ***
## X11          2.96699    0.06328  46.885 <2e-16 ***
## X12          1.91911    0.05636  34.050 <2e-16 ***
```

## X13	3.27055	0.06128	53.373	<2e-16	***
## X14	3.59485	0.06179	58.177	<2e-16	***
## X15	2.91142	0.06111	47.644	<2e-16	***
## X16	2.32401	0.06288	36.961	<2e-16	***
## X17	2.89123	0.05993	48.247	<2e-16	***
## X18	3.80344	0.06600	57.624	<2e-16	***
## X19	2.73710	0.06547	41.808	<2e-16	***
## X20	2.29227	0.06064	37.804	<2e-16	***
## X21	1.51705	0.06397	23.717	<2e-16	***
## X22	1.99664	0.06213	32.134	<2e-16	***
## X23	2.67939	0.06056	44.241	<2e-16	***
## X24	2.47020	0.06434	38.391	<2e-16	***
## X25	3.45721	0.06053	57.114	<2e-16	***
## X26	3.71335	0.06155	60.335	<2e-16	***
## X27	3.93587	0.07132	55.188	<2e-16	***
## X28	1.32561	0.06229	21.283	<2e-16	***
## X29	1.05586	0.06165	17.127	<2e-16	***
## X30	3.18332	0.05961	53.407	<2e-16	***
## X31	2.10471	0.06407	32.849	<2e-16	***
## X32	2.15854	0.06175	34.957	<2e-16	***
## X33	3.49628	0.06040	57.890	<2e-16	***
## X34	1.13456	0.06014	18.865	<2e-16	***
## X35	2.32519	0.07072	32.877	<2e-16	***
## X36	1.43736	0.05999	23.960	<2e-16	***
## X37	3.44222	0.06327	54.408	<2e-16	***
## X38	2.57043	0.06131	41.928	<2e-16	***
## X39	2.47086	0.06078	40.651	<2e-16	***
## X40	1.83546	0.05759	31.870	<2e-16	***
## X41	2.97487	0.06040	49.256	<2e-16	***
## X42	2.61514	0.07122	36.717	<2e-16	***
## X43	3.89452	0.06289	61.924	<2e-16	***
## X44	1.65122	0.05940	27.799	<2e-16	***
## X45	1.15451	0.06424	17.971	<2e-16	***
## X46	1.81052	0.06512	27.805	<2e-16	***
## X47	1.38634	0.06334	21.886	<2e-16	***
## X48	2.31795	0.06281	36.906	<2e-16	***
## X49	3.00236	0.06347	47.302	<2e-16	***
## X50	1.73813	0.06027	28.841	<2e-16	***
## X51	2.57682	0.05836	44.153	<2e-16	***
## X52	2.32508	0.06123	37.972	<2e-16	***
## X53	2.73581	0.06218	43.996	<2e-16	***
## X54	1.94659	0.06088	31.974	<2e-16	***
## X55	4.04786	0.05974	67.756	<2e-16	***
## X56	3.28758	0.06690	49.140	<2e-16	***
## X57	3.97055	0.06694	59.318	<2e-16	***
## X58	1.76909	0.05895	30.012	<2e-16	***
## X59	2.15947	0.06485	33.299	<2e-16	***
## X60	1.21264	0.06081	19.941	<2e-16	***
## X61	1.73856	0.06074	28.625	<2e-16	***
## X62	2.18242	0.06477	33.693	<2e-16	***
## X63	3.23326	0.06473	49.949	<2e-16	***
## X64	1.74698	0.06181	28.262	<2e-16	***
## X65	2.48855	0.06047	41.151	<2e-16	***
## X66	2.34540	0.06184	37.925	<2e-16	***


```
## X67      3.16775    0.06118   51.775   <2e-16 ***
## X68      2.07880    0.06094   34.115   <2e-16 ***
## X69      3.44976    0.06448   53.498   <2e-16 ***
## X70      2.28034    0.06063   37.611   <2e-16 ***
## X71      2.38348    0.06653   35.825   <2e-16 ***
## X72      3.96858    0.06071   65.365   <2e-16 ***
## X73      3.25039    0.06321   51.424   <2e-16 ***
## X74      2.73984    0.06052   45.274   <2e-16 ***
## X75      2.07016    0.06663   31.068   <2e-16 ***
## X76      2.17123    0.06721   32.304   <2e-16 ***
## X77      1.90020    0.06193   30.681   <2e-16 ***
## X78      3.48040    0.06083   57.211   <2e-16 ***
## X79      2.04801    0.06718   30.485   <2e-16 ***
## X80      2.37128    0.06071   39.061   <2e-16 ***
## X81      1.52090    0.06196   24.546   <2e-16 ***
## X82      1.39076    0.06557   21.211   <2e-16 ***
## X83      2.46487    0.06371   38.690   <2e-16 ***
## X84      2.47040    0.06129   40.309   <2e-16 ***
## X85      1.27911    0.06138   20.838   <2e-16 ***
## X86      1.98356    0.06161   32.193   <2e-16 ***
## X87      1.95006    0.05798   33.632   <2e-16 ***
## X88      3.11230    0.06466   48.135   <2e-16 ***
## X89      2.76429    0.06010   45.996   <2e-16 ***
## X90      1.51822    0.06968   21.789   <2e-16 ***
## X91      2.93123    0.06054   48.415   <2e-16 ***
## X92      1.71543    0.06225   27.555   <2e-16 ***
## X93      1.69781    0.05832   29.114   <2e-16 ***
## X94      1.72484    0.05832   29.576   <2e-16 ***
## X95      1.48116    0.06215   23.833   <2e-16 ***
## X96      2.89305    0.06567   44.052   <2e-16 ***
## X97      2.08861    0.06511   32.077   <2e-16 ***
## X98      2.91221    0.05995   48.581   <2e-16 ***
## X99      1.89725    0.06070   31.254   <2e-16 ***
## X100     1.47563    0.06213   23.752   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.49 on 155 degrees of freedom
## Multiple R-squared:  0.9994, Adjusted R-squared:  0.9991
## F-statistic: 2701 on 100 and 155 DF, p-value: < 2.2e-16
```

Your comment: We violated noise normality by replacing the normal distribution with a uniform distribution. Because the model had a relatively large dimension compared to the number of samples, the model started to overfit (nearly-perfect R^2 , very low RSE) when we replaced the normal-distributed noise with the uniform-distributed one.

```
noise <- X^4 %*% coefs

# KEEP THE CODE BELOW
Y <- (X %*% coefs) + intercept + noise
summary(learnAndTest(X, Y, "lm"))
```

Violate linearity

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-151753	-34912	959	34182	141637

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	487713.04	5359.69	90.997	< 2e-16 ***
X1	-96.29	1016.96	-0.095	0.92469
X2	1274.10	886.13	1.438	0.15250
X3	-244.52	942.08	-0.260	0.79555
X4	-1428.60	808.04	-1.768	0.07903 .
X5	-11.70	897.58	-0.013	0.98961
X6	433.64	899.01	0.482	0.63024
X7	-1006.95	932.03	-1.080	0.28164
X8	871.32	859.82	1.013	0.31246
X9	-1030.54	917.25	-1.124	0.26296
X10	-1323.38	954.59	-1.386	0.16764
X11	285.02	932.82	0.306	0.76036
X12	604.05	830.78	0.727	0.46827
X13	-1645.18	903.25	-1.821	0.07047 .
X14	-601.51	910.84	-0.660	0.50998
X15	-170.18	900.76	-0.189	0.85039
X16	111.43	926.84	0.120	0.90446
X17	-26.49	883.34	-0.030	0.97611
X18	-636.76	972.94	-0.654	0.51378
X19	1398.92	965.04	1.450	0.14919
X20	-560.23	893.80	-0.627	0.53172
X21	-1654.20	942.88	-1.754	0.08134 .
X22	20.03	915.89	0.022	0.98258
X23	-576.12	892.72	-0.645	0.51965
X24	874.30	948.44	0.922	0.35805
X25	1045.78	892.26	1.172	0.24297
X26	337.38	907.21	0.372	0.71048
X27	-27.84	1051.25	-0.026	0.97890
X28	-95.41	918.13	-0.104	0.91737
X29	841.02	908.75	0.925	0.35616
X30	1486.08	878.61	1.691	0.09277 .
X31	1304.68	944.46	1.381	0.16914
X32	-645.99	910.20	-0.710	0.47894
X33	252.33	890.26	0.283	0.77722
X34	-578.94	886.50	-0.653	0.51468
X35	-1267.75	1042.52	-1.216	0.22582
X36	-445.30	884.29	-0.504	0.61528
X37	-294.38	932.58	-0.316	0.75268
X38	-1027.43	903.68	-1.137	0.25732
X39	-353.10	895.95	-0.394	0.69405
X40	-639.90	848.93	-0.754	0.45213
X41	-240.88	890.27	-0.271	0.78708
X42	-175.78	1049.88	-0.167	0.86725
X43	-688.43	927.05	-0.743	0.45884

## X44	-988.75	875.56	-1.129	0.26052
## X45	-478.85	946.99	-0.506	0.61382
## X46	-462.79	959.83	-0.482	0.63037
## X47	889.38	933.72	0.953	0.34232
## X48	677.21	925.81	0.731	0.46559
## X49	-245.53	935.60	-0.262	0.79334
## X50	1443.80	888.35	1.625	0.10614
## X51	1508.28	860.26	1.753	0.08153 .
## X52	-565.08	902.59	-0.626	0.53219
## X53	-303.38	916.61	-0.331	0.74111
## X54	-1407.78	897.41	-1.569	0.11876
## X55	-1536.11	880.61	-1.744	0.08308 .
## X56	-1574.51	986.17	-1.597	0.11239
## X57	-1245.00	986.67	-1.262	0.20891
## X58	-130.82	868.90	-0.151	0.88052
## X59	-1778.56	955.92	-1.861	0.06470 .
## X60	-1343.72	896.40	-1.499	0.13590
## X61	1145.08	895.26	1.279	0.20279
## X62	492.03	954.78	0.515	0.60705
## X63	-285.58	954.17	-0.299	0.76511
## X64	1640.78	911.18	1.801	0.07369 .
## X65	-1005.53	891.40	-1.128	0.26105
## X66	-949.71	911.59	-1.042	0.29912
## X67	-233.11	901.86	-0.258	0.79638
## X68	624.73	898.21	0.696	0.48777
## X69	-1274.39	950.52	-1.341	0.18197
## X70	-584.19	893.70	-0.654	0.51428
## X71	-866.49	980.70	-0.884	0.37831
## X72	-334.03	894.95	-0.373	0.70948
## X73	-498.92	931.71	-0.535	0.59308
## X74	571.28	892.04	0.640	0.52285
## X75	-1745.80	982.19	-1.777	0.07745 .
## X76	-297.69	990.74	-0.300	0.76422
## X77	118.28	912.94	0.130	0.89708
## X78	-657.16	896.73	-0.733	0.46476
## X79	2198.82	990.29	2.220	0.02784 *
## X80	-712.94	894.84	-0.797	0.42683
## X81	-809.20	913.33	-0.886	0.37699
## X82	190.85	966.50	0.197	0.84372
## X83	-602.26	939.08	-0.641	0.52226
## X84	449.68	903.39	0.498	0.61935
## X85	956.20	904.82	1.057	0.29225
## X86	-1210.24	908.23	-1.333	0.18464
## X87	-511.84	854.68	-0.599	0.55013
## X88	44.99	953.08	0.047	0.96241
## X89	612.37	885.88	0.691	0.49044
## X90	-817.55	1027.10	-0.796	0.42726
## X91	2611.04	892.44	2.926	0.00395 **
## X92	857.08	917.67	0.934	0.35177
## X93	-1015.76	859.61	-1.182	0.23915
## X94	-139.75	859.65	-0.163	0.87107
## X95	1520.92	916.09	1.660	0.09889 .
## X96	1703.29	968.06	1.759	0.08047 .
## X97	241.73	959.79	0.252	0.80149

```
## X98          -1342.92      883.63  -1.520  0.13060
## X99          -1526.63      894.81  -1.706  0.08999 .
## X100          116.39       915.79   0.127  0.89903
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 66180 on 155 degrees of freedom
## Multiple R-squared:  0.4049, Adjusted R-squared:  0.02105
## F-statistic: 1.055 on 100 and 155 DF,  p-value: 0.3793
```

Your comment: We violated the linearity by adding a non-linear term (the fourth power, X^4 `%%` `coefs`). The F-statistic is almost equal to 1, which means that the model has no predictive capability (i.e., it is as good as an intercept-only model).

```
noise <- rnorm(n.samples, sd = seq(from = 4, to = 32, length.out = n.samples), mean = 0)

# KEEP THE CODE BELOW
Y <- (X %% coefs) + intercept + noise
summary(learnAndTest(X, Y, "lm"))
```

Violate homoscedasticity

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -45.860 -10.979  -0.396  11.107  40.398
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.3484     1.6338   0.213 0.831436
## X1             3.0153     0.3100   9.727 < 2e-16 ***
## X2             2.0975     0.2701   7.765 1.04e-12 ***
## X3             2.2721     0.2872   7.912 4.49e-13 ***
## X4             2.4563     0.2463   9.972 < 2e-16 ***
## X5             1.4821     0.2736   5.417 2.27e-07 ***
## X6             4.1105     0.2740  15.000 < 2e-16 ***
## X7             1.9203     0.2841   6.759 2.65e-10 ***
## X8             2.1342     0.2621   8.143 1.19e-13 ***
## X9             1.0557     0.2796   3.776 0.000227 ***
## X10            2.6741     0.2910   9.190 2.45e-16 ***
## X11            3.0767     0.2844  10.820 < 2e-16 ***
## X12            1.8402     0.2532   7.266 1.70e-11 ***
## X13            3.3236     0.2753  12.071 < 2e-16 ***
## X14            3.6122     0.2776  13.010 < 2e-16 ***
## X15            2.7170     0.2746   9.895 < 2e-16 ***
## X16            2.5454     0.2825   9.009 7.22e-16 ***
## X17            2.5162     0.2693   9.345 < 2e-16 ***
## X18            3.7894     0.2966  12.777 < 2e-16 ***
## X19            3.0005     0.2942  10.200 < 2e-16 ***
## X20            2.4748     0.2725   9.083 4.64e-16 ***
## X21            1.6190     0.2874   5.633 8.13e-08 ***
```

## X22	1.8612	0.2792	6.666	4.34e-10	***
## X23	2.8197	0.2721	10.362	< 2e-16	***
## X24	2.3143	0.2891	8.005	2.63e-13	***
## X25	3.6321	0.2720	13.354	< 2e-16	***
## X26	3.9784	0.2765	14.386	< 2e-16	***
## X27	4.1854	0.3205	13.061	< 2e-16	***
## X28	1.5043	0.2799	5.375	2.76e-07	***
## X29	0.9787	0.2770	3.533	0.000541	***
## X30	3.1974	0.2678	11.938	< 2e-16	***
## X31	2.2308	0.2879	7.748	1.14e-12	***
## X32	2.1960	0.2775	7.915	4.42e-13	***
## X33	3.3093	0.2714	12.194	< 2e-16	***
## X34	0.6634	0.2702	2.455	0.015190	*
## X35	2.4388	0.3178	7.674	1.74e-12	***
## X36	1.2266	0.2696	4.551	1.07e-05	***
## X37	3.5856	0.2843	12.613	< 2e-16	***
## X38	2.6125	0.2755	9.484	< 2e-16	***
## X39	2.0556	0.2731	7.526	3.99e-12	***
## X40	1.8130	0.2588	7.006	7.04e-11	***
## X41	3.2151	0.2714	11.847	< 2e-16	***
## X42	3.0515	0.3200	9.535	< 2e-16	***
## X43	4.1043	0.2826	14.524	< 2e-16	***
## X44	1.3685	0.2669	5.128	8.64e-07	***
## X45	0.9946	0.2887	3.445	0.000734	***
## X46	2.1493	0.2926	7.346	1.09e-11	***
## X47	1.5596	0.2846	5.480	1.69e-07	***
## X48	2.0556	0.2822	7.284	1.54e-11	***
## X49	2.7489	0.2852	9.638	< 2e-16	***
## X50	2.3142	0.2708	8.546	1.13e-14	***
## X51	2.3610	0.2622	9.004	7.48e-16	***
## X52	1.9467	0.2751	7.075	4.83e-11	***
## X53	2.7397	0.2794	9.805	< 2e-16	***
## X54	1.8172	0.2736	6.643	4.92e-10	***
## X55	4.2015	0.2684	15.652	< 2e-16	***
## X56	3.4543	0.3006	11.491	< 2e-16	***
## X57	3.6419	0.3008	12.109	< 2e-16	***
## X58	1.7907	0.2649	6.761	2.63e-10	***
## X59	2.6568	0.2914	9.117	3.78e-16	***
## X60	1.2353	0.2733	4.521	1.22e-05	***
## X61	1.6122	0.2729	5.908	2.12e-08	***
## X62	2.0774	0.2910	7.138	3.43e-11	***
## X63	2.8633	0.2909	9.844	< 2e-16	***
## X64	1.3366	0.2778	4.812	3.51e-06	***
## X65	2.4445	0.2717	8.996	7.81e-16	***
## X66	2.3827	0.2779	8.574	9.56e-15	***
## X67	3.5306	0.2749	12.843	< 2e-16	***
## X68	2.0217	0.2738	7.384	8.84e-12	***
## X69	3.5292	0.2897	12.180	< 2e-16	***
## X70	2.4073	0.2724	8.836	2.03e-15	***
## X71	2.9783	0.2989	9.962	< 2e-16	***
## X72	4.3121	0.2728	15.806	< 2e-16	***
## X73	3.3359	0.2840	11.745	< 2e-16	***
## X74	2.5283	0.2719	9.298	< 2e-16	***
## X75	2.0474	0.2994	6.838	1.74e-10	***

```
## X76          2.7044      0.3020      8.955 1.00e-15 ***
## X77          2.4624      0.2783      8.848 1.89e-15 ***
## X78          3.2735      0.2734     11.975 < 2e-16 ***
## X79          2.3339      0.3019      7.732 1.26e-12 ***
## X80          2.5439      0.2728      9.326 < 2e-16 ***
## X81          1.2723      0.2784      4.570 9.91e-06 ***
## X82          1.4678      0.2946      4.982 1.66e-06 ***
## X83          2.2225      0.2863      7.764 1.04e-12 ***
## X84          2.4679      0.2754      8.962 9.61e-16 ***
## X85          1.5728      0.2758      5.702 5.81e-08 ***
## X86          1.9681      0.2769      7.109 4.02e-11 ***
## X87          1.8789      0.2605      7.212 2.29e-11 ***
## X88          3.5052      0.2905     12.065 < 2e-16 ***
## X89          2.8293      0.2700     10.477 < 2e-16 ***
## X90          1.7530      0.3131      5.599 9.57e-08 ***
## X91          2.8495      0.2720     10.475 < 2e-16 ***
## X92          1.5911      0.2797      5.688 6.23e-08 ***
## X93          1.5468      0.2620      5.903 2.17e-08 ***
## X94          1.7570      0.2620      6.705 3.55e-10 ***
## X95          0.9938      0.2792      3.559 0.000495 ***
## X96          2.7859      0.2951      9.441 < 2e-16 ***
## X97          2.6419      0.2926      9.030 6.39e-16 ***
## X98          2.3278      0.2694      8.642 6.42e-15 ***
## X99          2.2463      0.2728      8.235 6.96e-14 ***
## X100         1.4841      0.2792      5.316 3.63e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 20.17 on 155 degrees of freedom
## Multiple R-squared:  0.9887, Adjusted R-squared:  0.9813
## F-statistic:   135 on 100 and 155 DF,  p-value: < 2.2e-16
```

Your comment: We violated the homoscedasticity by making the variance non-constant. Compared to the model trained on the original data, this model has much lower F-statistic and bigger RSE.

Understanding the advantages of shrinkage methods

In this part, we will be modifying the dependent variables X from task 1 by a linear transformation, represented by a square matrix of `n.dimensions` sides. For demonstration, consider that the identity function represented by the identity matrix:

```
M0 <- diag(n.dimensions) # Makes an identity matrix
X0 <- X %*% M0 # You can check that X0 == X
noise <- rnorm(n.samples, sd = 8, mean = 0)
Y0 <- (X0 %*% coefs) + noise # We can reuse the noise as it's independent of X.

# to see the original case
# print(learnAndTest(X0, Y0, "lm"))
# print(learnAndTest(X0, Y0, "ridge"))
# print(learnAndTest(X0, Y0, "lasso"))
```

Task 3

In this task, you will show your understanding of the advantages of Ridge by synthesizing data on which they perform the best. You are supposed to do this by linearly transforming the dataset X as in the example above.

In other words, you should construct a matrix which if used in place of M_0 in the above example would make Ridge perform better than the other two methods. You should not resort to degenerate cases where you would get a warning about using a rank-deficient matrix. Justify your method.

Scoring note: The difference in the RMSE criterion doesn't need to be large or does not need to be present if you are certain it's just a statistical artifact (which you could verify by re-running the tests multiple times or using the LOOCV in `learnAndTest` function). Your design and justification is what matters for the assignment evaluation and the measurements are here only to guide you. We expect this task to be challenging to students, but once you get the right idea, it is possible to implement it with very small amount of code.

```
# 1. Start with the identity transformation.
M1 <- diag(n.dimensions)
M1.fmin <- 0.2
M1.fmax <- 4
# 2. Introduce multicollinearity
#  $x_i = x_i + k_{i-1}x_{i-1} + k_{i-2}x_{i-2} + k_{i-3}x_{i-3}$ 
M1[row(M1) + 1 == col(M1)] <- seq(M1.fmin, M1.fmax, length = 99)
M1[row(M1) + 2 == col(M1)] <- seq(M1.fmax, M1.fmin, length = 98)
M1[row(M1) + 3 == col(M1)] <- seq(M1.fmin, M1.fmax, length = 97)
# do not change variable 1
M1[, 1] <- 0
M1[1, 1] <- 1

# KEEP THE CODE BELOW
noise <- rnorm(n.samples, sd = 8, mean = 0)
X1 <- X %*% M1
Y1 <- (X1 %*% coefs) + noise

# additional test code to see what M1 actually did
# X1.cov <- cov(X1)
# X1.cor <- round(cor(X1), 2)
# X1.cor_ <- X1.cor
# X1.cor_[X1.cor_ > 0.4] <- 1
# X1.cor_[X1.cor_ < 1] <- 0
# X1.cor__ <- colSums(X1.cor_)
# X1.pca <- prcomp(X1, scale. = TRUE)
# summary(X1.pca)
```

Ridge Running these test should now make Ridge perform the best.

```
print(learnAndTest(X1, Y1, "lm"))

## Linear Regression
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 231, 230, 230, 230, 232, 231, ...
## Resampling results:
##
##      RMSE      Rsquared    MAE
```

```

## 9.520415 0.9999154 7.543669
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
print(learnAndTest(X1, Y1, "ridge"))

## Ridge Regression
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 230, 232, 228, 231, 229, 232, ...
## Resampling results across tuning parameters:
##
##  lambda  RMSE          Rsquared  MAE
##  0e+00   1.537816e+09  0.6204258  1.136388e+09
##  1e-04   1.015722e+01  0.9999060  7.968957e+00
##  1e-01   6.507600e+01  0.9982386  5.249055e+01
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was lambda = 1e-04.
print(learnAndTest(X1, Y1, "lasso"))

## The lasso
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 232, 231, 231, 231, 232, 230, ...
## Resampling results across tuning parameters:
##
##  fraction  RMSE          Rsquared  MAE
##  0.1       8517813  0.6063046  6467118
##  0.5       42590048  0.6063004  32336324
##  0.9       76662282  0.6063001  58205530
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was fraction = 0.1.

```

Justification: Ridge works better than OLS when there is multicollinearity in the features (independent variables). We construct a transformation matrix such that it adds linear combination of variables x_{i-3} , x_{i-2} and x_{i-1} to every variable x_i (except x_1 , although it shouldn't matter). **We ran the test multiple times.** The way we transformed the data made LASSO perform visibly worse than Ridge and OLS (in terms of RMSE and R^2). Ridge and OLS gave similarly good results. However, **in most cases, Ridge performed best** (in terms of RMSE and R^2).

LASSO (OPTIONAL CHALLENGE) This part of the homework is purely optional, but we are eager to see students capable of solving this. Here you should do the same thing as above, but to make perform LASSO the best of the three methods. Although similar in nature, we consider this even more challenging than the above since being unable to modify the underlying coefficients, this may require some deeper considerations to justify the transformation method. It can still be implemented with a few short lines of

code, though.

```
# 1. Start with the identity transformation.
M2 <- diag(n.dimensions)
# 2. Let's randomly select 80 % of the features (which we'll make unimportant by scaling the values).
M2.unimportant_fraction <- 0.8
M2.selection <- sample(n.dimensions, replace = FALSE, size = round(M2.unimportant_fraction * n.dimensions))
M2.unimportant_idxes <- M2.selection
# 3. Update the transformation matrix so that it will scale down the selected features when applied.
# Note: We know the magnitude of the coefficients in our case, so we can choose correct scaling factor.
# We can also add the term `(1/coef[M2.unimportant_idxes])` (assuming abs(coef) > 1)
# to further normalize the scaling factor.
M2[, M2.unimportant_idxes] <- M2[, M2.unimportant_idxes] * 0.001 * (1 / coef[M2.unimportant_idxes])
# M2[,M2.unimportant_idxes] <- M2[,M2.unimportant_idxes] * 0.001 # also works great

# KEEP THE CODE BELOW
noise <- rnorm(n.samples, sd = 8, mean = 0)
X2 <- X %*% M2
Y2 <- (X2 %*% coef) + noise
```

Running these test should now make LASSO perform the best.

```
print(learnAndTest(X2, Y2, "lm"))
```

```
## Linear Regression
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 230, 228, 230, 231, 231, 231, ...
## Resampling results:
##
##      RMSE      Rsquared   MAE
##  10.51489   0.9738668   8.722386
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
print(learnAndTest(X2, Y2, "ridge"))
```

```
## Ridge Regression
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 231, 231, 231, 230, 231, 232, ...
## Resampling results across tuning parameters:
##
##   lambda  RMSE      Rsquared   MAE
##   0e+00   10.48616  0.9740718   8.461986
##   1e-04   10.48573  0.9740723   8.460979
##   1e-01   12.73543  0.9629104  10.092765
##
```

```
## RMSE was used to select the optimal model using the smallest value.  
## The final value used for the model was lambda = 1e-04.
```

```
print(learnAndTest(X2, Y2, "lasso"))
```

```
## The lasso  
##  
## 256 samples  
## 100 predictors  
##  
## No pre-processing  
## Resampling: Cross-Validated (10 fold)  
## Summary of sample sizes: 230, 230, 229, 231, 231, ...  
## Resampling results across tuning parameters:  
##  
##   fraction  RMSE      Rsquared   MAE  
##   0.1       55.990259  0.5183220  45.212815  
##   0.5       26.994841  0.9195134  21.729017  
##   0.9        9.185182  0.9802712   7.298981  
##  
## RMSE was used to select the optimal model using the smallest value.  
## The final value used for the model was fraction = 0.9.
```

Justification: Generally, LASSO performs better when the response is a function of only a relatively small number of predictors. We can adjust the data by a linear transformation to make it the case. We create a scaling matrix that scales down a selected subset of features (making them unimportant) which also affects the linear regression coefficients. This in turn pushes LASSO to perform variable selection (forcing some of the coefficient estimates to be exactly equal to zero when lambda is sufficiently large). We ran the tests multiple times for our transformation and in all instances, the LASSO was the best (lowest RMSE, with a difference of around 1).