

SAN Assignment - regression

Martin Endler endlemar@fel.cvut.cz

Submission

Fill in your name above for clarity. To solve this homework, simply write your answers into this document and fill in the marked pieces of code. Submit your solution consisting of both this modified Rmd file and a knitted PDF document as an archive to the courseware BRUTE upload system for the SAN course. The deadline is specified there.

Initialization

Load the required libraries `gtools`, `caret` and `glmnet`, make sure you have those installed. We also fix the random seed for reproducibility convenience.

```
require(gtools);

## Loading required package: gtools
require(caret);

## Loading required package: caret
## Loading required package: ggplot2
## Loading required package: lattice
require(glmnet);

## Loading required package: glmnet
## Loading required package: Matrix
## Loaded glmnet 4.1-4
##
## Attaching package: 'glmnet'
## The following object is masked from 'package:gtools':
##
##   na.replace
set.seed(0)
```

Here, we define constants of the assignment. You may play with the values and observe what happens, but in your solution you should use the given values unchanged.

```
n.samples <- 256 # Total number of samples (training and testing together)
n.dimensions <- 100 # Number of n.dimensions, a.k.a. attributes or features
```

Model evaluation procedure

The function `learnAndTest` takes a matrix of independent variables values `X` and a corresponding vector of dependent variable (a.k.a. response) `Y` then trains and evaluates a model specified by the `modelType` parameter.

If you are interested in all possible `modelType` parameter values, refer to <http://topepo.github.io/caret/train-models-by-tag.html>. Here we will only be using three: "lm" for ordinary least squares and "glmnet" with parameter `alpha = 1` resp. `alpha = 0` for LASSO resp. Ridge. (There are also `lasso` and `ridge` methods you may try, but these have inconsistent API for passing `lambda`.) This function learns the model from the data and estimates its accuracy using cross validation. The results are then printed to output. For purpose of this assignment, consider only the RMSE (root-mean-square error) criterion.

```
learnAndTest <- function(X, Y, modelType, ...) {
  data <- data.frame(X, Y)

  train.control <- trainControl(method = "cv", number = 10) # 10-fold cross-validation
  # Alternative to get more accurate, but slower evaluation:
  # train.control <- trainControl(method = "LOOCV")

  # Here we train the model using the versatile `train` function from the caret package
  train(Y ~ .,
        data = data,
        method = modelType,
        trControl = train.control,
        ...
  )
}
```

We will also precompute an array of candidate `lambda` values for LASSO and Ridge.

```
lambda_lasso <- expand.grid(lambda = 10^seq(10, -3, length = 10), alpha = 1)
lambda_ridge <- expand.grid(lambda = 10^seq(10, -3, length = 10), alpha = 0)
```

Initial data generation

Here, we generate some data.

```
# Generates independent variables by uniform i.i.d. sampling
X <- replicate(
  n.dimensions,
  runif(n.samples, min = -10, max = 10)
)

# Randomly generates the actual underlying coefficients of the linear dependency
coefs <- runif(n.dimensions, min = 1, max = 4)
intercept <- 0 # For simplicity

# Synthesizes dependent variable (observed values) by the given linear dependency plus noise
noise <- rnorm(n.samples, sd = 8, mean = 0) # Gaussian noise to be added to the response
Y <- (X %*% coefs) + intercept + noise # Note: (%*%) is the matrix multiplication operator
```

Testing the models

Now let us run the following tests:

```
print(learnAndTest(X, Y, "lm"))
```

```
## Linear Regression
##
## 256 samples
## 100 predictors
```

```

##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 231, 231, 230, 231, 231, 229, ...
## Resampling results:
##
##      RMSE      Rsquared    MAE
##      10.22282  0.9954098  8.107877
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
print(learnAndTest(X, Y, "glmnet", tuneGrid = lambda_ridge))

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## glmnet
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 229, 231, 231, 231, 228, 230, ...
## Resampling results across tuning parameters:
##
##      lambda      RMSE      Rsquared    MAE
##      1.000000e-03  13.58962  0.9933283  11.09131
##      2.782559e-02  13.58962  0.9933283  11.09131
##      7.742637e-01  13.58962  0.9933283  11.09131
##      2.154435e+01  34.01922  0.9740668  27.24431
##      5.994843e+02  120.12106  0.7891572  97.88858
##      1.668101e+04  143.78427  0.7190257  117.35115
##      4.641589e+05  144.96674      NaN  118.32455
##      1.291550e+07  144.96674      NaN  118.32455
##      3.593814e+08  144.96674      NaN  118.32455
##      1.000000e+10  144.96674      NaN  118.32455
##
## Tuning parameter 'alpha' was held constant at a value of 0
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 0 and lambda = 0.7742637.
print(learnAndTest(X, Y, "glmnet", tuneGrid = lambda_lasso))

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## glmnet
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 230, 230, 230, 231, 228, 232, ...
## Resampling results across tuning parameters:

```

```
##
##      lambda      RMSE      Rsquared      MAE
##  1.000000e-03   10.43912  0.9950887    8.530893
##  2.782559e-02   10.43912  0.9950887    8.530893
##  7.742637e-01   21.50918  0.9851375   16.937733
##  2.154435e+01  135.67380  0.1978032  109.662368
##  5.994843e+02  145.51104      NaN   118.523560
##  1.668101e+04  145.51104      NaN   118.523560
##  4.641589e+05  145.51104      NaN   118.523560
##  1.291550e+07  145.51104      NaN   118.523560
##  3.593814e+08  145.51104      NaN   118.523560
##  1.000000e+10  145.51104      NaN   118.523560
##
## Tuning parameter 'alpha' was held constant at a value of 1
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 1 and lambda = 0.02782559.
```

Task 1:

Answer the following questions:

- What change in the learned model would you anticipate if we changed the **mean** parameter value to a different constant in the noise generation? You may answer this either by talking about the coefficients or by giving a geometrical interpretation. **Answer:** From a geometric point of view, the predicted y values would be shifted up (for a positive mean), resp. down (for a negative mean). This shift would correspond to a change in the estimated intercept parameter. The other estimated coefficients would remain the same.
- In our example we generated samples (i.e. the independent variables X) from uniform distribution. The least squares method, on the other hand, has something called the “normality assumption”. Have we violated that assumption? Justify. **Answer:** No, we haven’t. The “normality assumption” requires that the residuals are normally distributed, not the data (our generated samples).
- Which method gave the best results? Is it the most common one to do so if you re-run the test several times? Why do you think it performs better the best? **Answer:** The LASSO gave the best result (smallest RMSE) most often while re-running the test multiple times. However, the plain Least squares method was also very good, with a small RMSE very close to the one from LASSO. From the very small lambda parameter in the best LASSO, we can see that the LASSO is almost the same as the Least squares method (the shrinkage penalty term is close to zero). This could mean that the Least squares method is not overfitting on our data.
- Check the selected values for the lambda parameter for ridge and LASSO. Are they low or high? How does it relate to the above answer? **Answer:** They are both very low. That means that both ridge and LASSO are almost the same as the Least squares method (the shrinkage penalty term is close to zero).

Least squares assumptions

The data generation model assumed by the ordinary least squares method can be mathematically written as follows:

$$Y = \mathbf{X}^T \boldsymbol{\beta} + \beta_0 + G, \quad G \sim \mathcal{N}(0, \sigma^2)$$

This formula implicitly expresses some of the assumptions about the data, required for the method to work reliably.

- The observed value Y is influenced by some Gaussian noise G .
- There truly exists an underlying linear dependency.
- The noise is homoscedastic (σ^2 is a constant).

Task 2:

First of all, make sure you understand how elements of this formula correspond to the code in the “data generation” section.

Your task is to violate each of these assumptions (one at a time) and **briefly** comment the changes in the learned model by statistically comparing it to model using the above data. (Coefficient summary below.) The catch here is that you are allowed to only modify the noise generation procedure to achieve that. Attempt to find a way of violating the assumptions to achieve a clear difference, but any solution that is technically correct will be awarded full points.

It is sufficient to look at the **summary** of the OLS model.

```
summary(learnAndTest(X, Y, "lm"))
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-16.3220	-3.7072	0.1499	4.0184	16.0036

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.43023	0.63316	-0.679	0.498
X1	3.32876	0.12014	27.708	<2e-16 ***
X2	1.61195	0.10468	15.399	<2e-16 ***
X3	2.42437	0.11129	21.784	<2e-16 ***
X4	2.90811	0.09546	30.466	<2e-16 ***
X5	1.15971	0.10603	10.937	<2e-16 ***
X6	3.82874	0.10620	36.051	<2e-16 ***
X7	2.37577	0.11010	21.578	<2e-16 ***
X8	1.94196	0.10157	19.119	<2e-16 ***
X9	1.26780	0.10836	11.700	<2e-16 ***
X10	2.43546	0.11277	21.597	<2e-16 ***
X11	2.79039	0.11020	25.322	<2e-16 ***
X12	1.84899	0.09814	18.840	<2e-16 ***
X13	3.31907	0.10670	31.105	<2e-16 ***
X14	3.73201	0.10760	34.684	<2e-16 ***
X15	2.85848	0.10641	26.863	<2e-16 ***
X16	2.18332	0.10949	19.941	<2e-16 ***
X17	2.98273	0.10435	28.583	<2e-16 ***
X18	3.95243	0.11494	34.388	<2e-16 ***
X19	2.72129	0.11400	23.870	<2e-16 ***
X20	2.26276	0.10559	21.430	<2e-16 ***
X21	1.45329	0.11139	13.047	<2e-16 ***
X22	2.06841	0.10820	19.117	<2e-16 ***
X23	2.65462	0.10546	25.172	<2e-16 ***
X24	2.43848	0.11204	21.764	<2e-16 ***
X25	3.39725	0.10541	32.230	<2e-16 ***
X26	3.63669	0.10717	33.933	<2e-16 ***
X27	4.08163	0.12419	32.867	<2e-16 ***
X28	1.39751	0.10846	12.885	<2e-16 ***
X29	1.21262	0.10735	11.296	<2e-16 ***
X30	3.36342	0.10379	32.405	<2e-16 ***

## X31	1.98639	0.11157	17.804	<2e-16	***
## X32	2.06385	0.10752	19.194	<2e-16	***
## X33	3.41435	0.10517	32.465	<2e-16	***
## X34	1.04075	0.10473	9.938	<2e-16	***
## X35	2.57185	0.12316	20.883	<2e-16	***
## X36	1.44023	0.10446	13.787	<2e-16	***
## X37	3.28405	0.11017	29.809	<2e-16	***
## X38	2.72006	0.10675	25.480	<2e-16	***
## X39	2.26690	0.10584	21.418	<2e-16	***
## X40	1.88773	0.10029	18.823	<2e-16	***
## X41	2.95467	0.10517	28.094	<2e-16	***
## X42	2.47410	0.12403	19.948	<2e-16	***
## X43	3.85890	0.10951	35.236	<2e-16	***
## X44	1.37370	0.10343	13.281	<2e-16	***
## X45	1.31642	0.11187	11.767	<2e-16	***
## X46	1.89909	0.11339	16.749	<2e-16	***
## X47	1.35180	0.11030	12.255	<2e-16	***
## X48	2.36427	0.10937	21.617	<2e-16	***
## X49	3.06998	0.11053	27.776	<2e-16	***
## X50	1.69993	0.10494	16.199	<2e-16	***
## X51	2.66625	0.10163	26.236	<2e-16	***
## X52	2.27966	0.10663	21.380	<2e-16	***
## X53	2.46486	0.10828	22.763	<2e-16	***
## X54	2.02121	0.10601	19.065	<2e-16	***
## X55	4.06324	0.10403	39.058	<2e-16	***
## X56	3.48045	0.11650	29.875	<2e-16	***
## X57	3.82223	0.11656	32.792	<2e-16	***
## X58	1.83125	0.10265	17.841	<2e-16	***
## X59	2.07426	0.11293	18.368	<2e-16	***
## X60	1.01282	0.10589	9.564	<2e-16	***
## X61	1.88231	0.10576	17.798	<2e-16	***
## X62	2.19831	0.11279	19.490	<2e-16	***
## X63	3.36844	0.11272	29.884	<2e-16	***
## X64	1.62348	0.10764	15.082	<2e-16	***
## X65	2.63364	0.10530	25.010	<2e-16	***
## X66	2.26839	0.10769	21.064	<2e-16	***
## X67	3.15058	0.10654	29.572	<2e-16	***
## X68	2.11758	0.10611	19.957	<2e-16	***
## X69	3.48378	0.11229	31.026	<2e-16	***
## X70	2.20811	0.10557	20.915	<2e-16	***
## X71	2.45733	0.11585	21.211	<2e-16	***
## X72	3.93237	0.10572	37.195	<2e-16	***
## X73	3.03948	0.11007	27.615	<2e-16	***
## X74	2.77817	0.10538	26.363	<2e-16	***
## X75	2.04650	0.11603	17.638	<2e-16	***
## X76	2.18996	0.11704	18.711	<2e-16	***
## X77	1.84367	0.10785	17.095	<2e-16	***
## X78	3.64003	0.10593	34.361	<2e-16	***
## X79	2.09904	0.11699	17.943	<2e-16	***
## X80	2.31115	0.10571	21.863	<2e-16	***
## X81	1.60042	0.10789	14.833	<2e-16	***
## X82	1.34889	0.11418	11.814	<2e-16	***
## X83	2.51265	0.11094	22.649	<2e-16	***
## X84	2.70763	0.10672	25.371	<2e-16	***

```
## X85      1.17989    0.10689   11.038   <2e-16 ***
## X86      2.10964    0.10729   19.663   <2e-16 ***
## X87      2.16377    0.10097   21.431   <2e-16 ***
## X88      3.07060    0.11259   27.272   <2e-16 ***
## X89      2.72039    0.10465   25.995   <2e-16 ***
## X90      1.51493    0.12133   12.486   <2e-16 ***
## X91      2.63591    0.10543   25.002   <2e-16 ***
## X92      1.78357    0.10841   16.453   <2e-16 ***
## X93      1.66301    0.10155   16.377   <2e-16 ***
## X94      1.56552    0.10155   15.416   <2e-16 ***
## X95      1.47395    0.10822   13.620   <2e-16 ***
## X96      2.82689    0.11436   24.719   <2e-16 ***
## X97      2.07620    0.11338   18.311   <2e-16 ***
## X98      2.81395    0.10439   26.957   <2e-16 ***
## X99      1.97136    0.10571   18.649   <2e-16 ***
## X100     1.65748    0.10818   15.321   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.818 on 155 degrees of freedom
## Multiple R-squared:  0.9983, Adjusted R-squared:  0.9971
## F-statistic: 888.8 on 100 and 155 DF,  p-value: < 2.2e-16
```


```
noise2 <- rbeta(n.samples, 1 / 7, 1 / 7)
# noise3 <- ((noise2 - mean(noise2))/sd(noise2) * 8)
noise3 <- ((seq(n.samples) - mean(seq(n.samples))) / sd(seq(n.samples))) * 8

noise2 <- runif(n.samples, min = -8, max = 8)
noise <- rnorm(n.samples, sd = 8, mean = 0)
mean(noise3)
```


Violate noise normality

```
## [1] 0
var(noise3)


## [1] 64
plot(noise)
```




```
plot(noise3)
```

```
D3 <- (X %*% coefs) + intercept + noise3  
D1 <- (X %*% coefs) + intercept + noise  
plot(D1)
```



```
plot(D3)
```



```
# KEEP THE CODE BELOW
Y <- (X %*% coefs) + intercept + noise3
# summary(learnAndTest(X, Y, "lm"))
test_model <- lm(Y ~ X)
summary(test_model)
```

```
##
## Call:
## lm(formula = Y ~ X)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-15.8447	-4.6926	0.0893	4.4805	17.9361

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.18119	0.66154	-0.274	0.785
X1	3.30541	0.12552	26.333	<2e-16 ***
X2	1.84935	0.10937	16.909	<2e-16 ***
X3	2.26219	0.11628	19.455	<2e-16 ***
X4	2.77587	0.09974	27.832	<2e-16 ***
X5	1.08395	0.11079	9.784	<2e-16 ***
X6	3.83938	0.11096	34.600	<2e-16 ***
X7	2.06626	0.11504	17.961	<2e-16 ***
X8	1.89383	0.10613	17.845	<2e-16 ***
X9	1.05184	0.11321	9.291	<2e-16 ***

## X10	2.55281	0.11782	21.666	<2e-16	***
## X11	2.80752	0.11514	24.384	<2e-16	***
## X12	1.74527	0.10254	17.020	<2e-16	***
## X13	3.31791	0.11149	29.760	<2e-16	***
## X14	3.77614	0.11242	33.589	<2e-16	***
## X15	2.81804	0.11118	25.347	<2e-16	***
## X16	2.44700	0.11440	21.390	<2e-16	***
## X17	2.89316	0.10903	26.536	<2e-16	***
## X18	3.74582	0.12009	31.192	<2e-16	***
## X19	2.79024	0.11911	23.425	<2e-16	***
## X20	2.16422	0.11032	19.618	<2e-16	***
## X21	1.47295	0.11638	12.656	<2e-16	***
## X22	2.03095	0.11305	17.966	<2e-16	***
## X23	2.53931	0.11019	23.045	<2e-16	***
## X24	2.58547	0.11706	22.086	<2e-16	***
## X25	3.66204	0.11013	33.252	<2e-16	***
## X26	3.76359	0.11198	33.611	<2e-16	***
## X27	4.12068	0.12975	31.757	<2e-16	***
## X28	1.43790	0.11332	12.688	<2e-16	***
## X29	1.21979	0.11217	10.875	<2e-16	***
## X30	3.41967	0.10845	31.534	<2e-16	***
## X31	2.00793	0.11657	17.225	<2e-16	***
## X32	2.16965	0.11235	19.312	<2e-16	***
## X33	3.54981	0.10988	32.305	<2e-16	***
## X34	1.19071	0.10942	10.882	<2e-16	***
## X35	2.54635	0.12868	19.789	<2e-16	***
## X36	1.29689	0.10915	11.882	<2e-16	***
## X37	3.43179	0.11511	29.814	<2e-16	***
## X38	2.63968	0.11154	23.666	<2e-16	***
## X39	2.28773	0.11059	20.687	<2e-16	***
## X40	1.95932	0.10478	18.699	<2e-16	***
## X41	3.24824	0.10988	29.560	<2e-16	***
## X42	2.65336	0.12959	20.476	<2e-16	***
## X43	3.85008	0.11442	33.647	<2e-16	***
## X44	1.56029	0.10807	14.438	<2e-16	***
## X45	1.18534	0.11689	10.141	<2e-16	***
## X46	1.81464	0.11847	15.317	<2e-16	***
## X47	1.33129	0.11525	11.552	<2e-16	***
## X48	2.27166	0.11427	19.879	<2e-16	***
## X49	3.01522	0.11548	26.110	<2e-16	***
## X50	1.66570	0.10965	15.191	<2e-16	***
## X51	2.65017	0.10618	24.959	<2e-16	***
## X52	2.50679	0.11141	22.502	<2e-16	***
## X53	2.68830	0.11314	23.762	<2e-16	***
## X54	2.09325	0.11077	18.898	<2e-16	***
## X55	4.16219	0.10869	38.293	<2e-16	***
## X56	3.37611	0.12172	27.736	<2e-16	***
## X57	3.90781	0.12178	32.088	<2e-16	***
## X58	1.81200	0.10725	16.896	<2e-16	***
## X59	2.19546	0.11799	18.607	<2e-16	***
## X60	1.02624	0.11064	9.275	<2e-16	***
## X61	1.69194	0.11050	15.311	<2e-16	***
## X62	2.24581	0.11785	19.057	<2e-16	***
## X63	3.22959	0.11777	27.422	<2e-16	***

```
## X64      1.67417    0.11247   14.886   <2e-16 ***
## X65      2.43306    0.11002   22.114   <2e-16 ***
## X66      2.31424    0.11252   20.568   <2e-16 ***
## X67      3.36971    0.11132   30.272   <2e-16 ***
## X68      1.86476    0.11087   16.820   <2e-16 ***
## X69      3.41243    0.11732   29.086   <2e-16 ***
## X70      2.25226    0.11031   20.418   <2e-16 ***
## X71      2.37516    0.12105   19.622   <2e-16 ***
## X72      3.90285    0.11046   35.332   <2e-16 ***
## X73      3.32479    0.11500   28.911   <2e-16 ***
## X74      2.77866    0.11010   25.237   <2e-16 ***
## X75      2.12503    0.12123   17.529   <2e-16 ***
## X76      2.27454    0.12229   18.600   <2e-16 ***
## X77      1.82027    0.11268   16.154   <2e-16 ***
## X78      3.55086    0.11068   32.081   <2e-16 ***
## X79      2.25025    0.12223   18.410   <2e-16 ***
## X80      2.38035    0.11045   21.552   <2e-16 ***
## X81      1.57821    0.11273   14.000   <2e-16 ***
## X82      1.31645    0.11929   11.035   <2e-16 ***
## X83      2.57145    0.11591   22.185   <2e-16 ***
## X84      2.44245    0.11150   21.905   <2e-16 ***
## X85      1.14263    0.11168   10.231   <2e-16 ***
## X86      1.74402    0.11210   15.557   <2e-16 ***
## X87      1.98537    0.10549   18.820   <2e-16 ***
## X88      3.16381    0.11764   26.895   <2e-16 ***
## X89      2.70078    0.10934   24.700   <2e-16 ***
## X90      1.51776    0.12677   11.972   <2e-16 ***
## X91      2.69703    0.11015   24.484   <2e-16 ***
## X92      1.52576    0.11327   13.471   <2e-16 ***
## X93      1.82118    0.10610   17.165   <2e-16 ***
## X94      1.50756    0.10611   14.208   <2e-16 ***
## X95      1.47727    0.11307   13.065   <2e-16 ***
## X96      2.96730    0.11949   24.834   <2e-16 ***
## X97      2.34247    0.11847   19.773   <2e-16 ***
## X98      2.76990    0.10907   25.397   <2e-16 ***
## X99      2.00603    0.11044   18.163   <2e-16 ***
## X100     1.33923    0.11304   11.848   <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 8.169 on 155 degrees of freedom
```

```
## Multiple R-squared:  0.9981, Adjusted R-squared:  0.9969
```

```
## F-statistic: 816.9 on 100 and 155 DF,  p-value: < 2.2e-16
```

```
# plot(test_model)
```

Your comment: We violated noise normality by replacing the normal distribution with uniform distribution (with the 0 mean and

```
noise <- X^4 %%% coefs
```

```
# KEEP THE CODE BELOW
```

```
Y <- (X %%% coefs) + intercept + noise
```

```
summary(learnAndTest(X, Y, "lm"))
```

Violate linearity

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -151753  -34912    959    34182   141637
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 487713.04    5359.69   90.997 < 2e-16 ***
## X1           -96.29     1016.96   -0.095  0.92469
## X2            1274.10      886.13    1.438  0.15250
## X3           -244.52      942.08   -0.260  0.79555
## X4          -1428.60      808.04   -1.768  0.07903 .
## X5            -11.70      897.58   -0.013  0.98961
## X6             433.64      899.01    0.482  0.63024
## X7          -1006.95      932.03   -1.080  0.28164
## X8             871.32      859.82    1.013  0.31246
## X9          -1030.54      917.25   -1.124  0.26296
## X10          -1323.38      954.59   -1.386  0.16764
## X11             285.02      932.82    0.306  0.76036
## X12             604.05      830.78    0.727  0.46827
## X13          -1645.18      903.25   -1.821  0.07047 .
## X14           -601.51      910.84   -0.660  0.50998
## X15          -170.18      900.76   -0.189  0.85039
## X16             111.43      926.84    0.120  0.90446
## X17            -26.49      883.34   -0.030  0.97611
## X18           -636.76      972.94   -0.654  0.51378
## X19           1398.92      965.04    1.450  0.14919
## X20          -560.23      893.80   -0.627  0.53172
## X21          -1654.20      942.88   -1.754  0.08134 .
## X22             20.03      915.89    0.022  0.98258
## X23          -576.12      892.72   -0.645  0.51965
## X24             874.30      948.44    0.922  0.35805
## X25           1045.78      892.26    1.172  0.24297
## X26             337.38      907.21    0.372  0.71048
## X27            -27.84     1051.25   -0.026  0.97890
## X28            -95.41      918.13   -0.104  0.91737
## X29             841.02      908.75    0.925  0.35616
## X30           1486.08      878.61    1.691  0.09277 .
## X31           1304.68      944.46    1.381  0.16914
## X32          -645.99      910.20   -0.710  0.47894
## X33             252.33      890.26    0.283  0.77722
## X34          -578.94      886.50   -0.653  0.51468
## X35          -1267.75     1042.52   -1.216  0.22582
## X36          -445.30      884.29   -0.504  0.61528
## X37          -294.38      932.58   -0.316  0.75268
## X38          -1027.43      903.68   -1.137  0.25732
## X39          -353.10      895.95   -0.394  0.69405
```

## X40	-639.90	848.93	-0.754	0.45213
## X41	-240.88	890.27	-0.271	0.78708
## X42	-175.78	1049.88	-0.167	0.86725
## X43	-688.43	927.05	-0.743	0.45884
## X44	-988.75	875.56	-1.129	0.26052
## X45	-478.85	946.99	-0.506	0.61382
## X46	-462.79	959.83	-0.482	0.63037
## X47	889.38	933.72	0.953	0.34232
## X48	677.21	925.81	0.731	0.46559
## X49	-245.53	935.60	-0.262	0.79334
## X50	1443.80	888.35	1.625	0.10614
## X51	1508.28	860.26	1.753	0.08153 .
## X52	-565.08	902.59	-0.626	0.53219
## X53	-303.38	916.61	-0.331	0.74111
## X54	-1407.78	897.41	-1.569	0.11876
## X55	-1536.11	880.61	-1.744	0.08308 .
## X56	-1574.51	986.17	-1.597	0.11239
## X57	-1245.00	986.67	-1.262	0.20891
## X58	-130.82	868.90	-0.151	0.88052
## X59	-1778.56	955.92	-1.861	0.06470 .
## X60	-1343.72	896.40	-1.499	0.13590
## X61	1145.08	895.26	1.279	0.20279
## X62	492.03	954.78	0.515	0.60705
## X63	-285.58	954.17	-0.299	0.76511
## X64	1640.78	911.18	1.801	0.07369 .
## X65	-1005.53	891.40	-1.128	0.26105
## X66	-949.71	911.59	-1.042	0.29912
## X67	-233.11	901.86	-0.258	0.79638
## X68	624.73	898.21	0.696	0.48777
## X69	-1274.39	950.52	-1.341	0.18197
## X70	-584.19	893.70	-0.654	0.51428
## X71	-866.49	980.70	-0.884	0.37831
## X72	-334.03	894.95	-0.373	0.70948
## X73	-498.92	931.71	-0.535	0.59308
## X74	571.28	892.04	0.640	0.52285
## X75	-1745.80	982.19	-1.777	0.07745 .
## X76	-297.69	990.74	-0.300	0.76422
## X77	118.28	912.94	0.130	0.89708
## X78	-657.16	896.73	-0.733	0.46476
## X79	2198.82	990.29	2.220	0.02784 *
## X80	-712.94	894.84	-0.797	0.42683
## X81	-809.20	913.33	-0.886	0.37699
## X82	190.85	966.50	0.197	0.84372
## X83	-602.26	939.08	-0.641	0.52226
## X84	449.68	903.39	0.498	0.61935
## X85	956.20	904.82	1.057	0.29225
## X86	-1210.24	908.23	-1.333	0.18464
## X87	-511.84	854.68	-0.599	0.55013
## X88	44.99	953.08	0.047	0.96241
## X89	612.37	885.88	0.691	0.49044
## X90	-817.55	1027.10	-0.796	0.42726
## X91	2611.04	892.44	2.926	0.00395 **
## X92	857.08	917.67	0.934	0.35177
## X93	-1015.76	859.61	-1.182	0.23915

```
## X94          -139.75      859.65  -0.163  0.87107
## X95          1520.92     916.09   1.660  0.09889 .
## X96          1703.29     968.06   1.759  0.08047 .
## X97           241.73     959.79   0.252  0.80149
## X98         -1342.92     883.63  -1.520  0.13060
## X99         -1526.63     894.81  -1.706  0.08999 .
## X100          116.39     915.79   0.127  0.89903
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 66180 on 155 degrees of freedom
## Multiple R-squared:  0.4049, Adjusted R-squared:  0.02105
## F-statistic: 1.055 on 100 and 155 DF,  p-value: 0.3793
```

Your comment: We violated the linearity by adding a non-linear term (the fourth power, X^4 `coefs`). The F-statistic is almost equal to 1, which means that the model has no predictive capability (i.e., it is as good as an intercept-only model).

```
noise <- rnorm(n.samples, sd = seq(from = 4, to = 32, length.out = n.samples), mean = 0)

# KEEP THE CODE BELOW
Y <- (X %*% coefs) + intercept + noise
summary(learnAndTest(X, Y, "lm"))
```

Violate homoscedasticity

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -51.091 -10.739   1.186  10.119  50.303
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.8570     1.7884   1.038 0.300700
## X1             2.8583     0.3393   8.423 2.32e-14 ***
## X2             1.2512     0.2957   4.232 3.96e-05 ***
## X3             2.2611     0.3143   7.193 2.54e-11 ***
## X4             3.0562     0.2696  11.335 < 2e-16 ***
## X5             1.4764     0.2995   4.929 2.10e-06 ***
## X6             3.5592     0.3000  11.865 < 2e-16 ***
## X7             2.4826     0.3110   7.983 2.99e-13 ***
## X8             2.2331     0.2869   7.784 9.33e-13 ***
## X9             1.3874     0.3061   4.533 1.16e-05 ***
## X10            2.5504     0.3185   8.007 2.60e-13 ***
## X11            3.0665     0.3113   9.852 < 2e-16 ***
## X12            1.7705     0.2772   6.387 1.88e-09 ***
## X13            3.1827     0.3014  10.560 < 2e-16 ***
## X14            3.9118     0.3039  12.871 < 2e-16 ***
## X15            3.2312     0.3006  10.751 < 2e-16 ***
## X16            2.6005     0.3093   8.409 2.53e-14 ***
## X17            2.6135     0.2947   8.867 1.69e-15 ***
```


## X18	4.3989	0.3246	13.550	< 2e-16	***
## X19	2.6628	0.3220	8.269	5.71e-14	***
## X20	2.4564	0.2982	8.236	6.91e-14	***
## X21	1.3050	0.3146	4.148	5.51e-05	***
## X22	1.9173	0.3056	6.274	3.36e-09	***
## X23	2.5237	0.2979	8.472	1.74e-14	***
## X24	2.1802	0.3165	6.889	1.32e-10	***
## X25	3.3679	0.2977	11.312	< 2e-16	***
## X26	3.7960	0.3027	12.540	< 2e-16	***
## X27	3.9518	0.3508	11.266	< 2e-16	***
## X28	1.0275	0.3064	3.354	0.001001	**
## X29	0.9684	0.3032	3.194	0.001702	**
## X30	3.0893	0.2932	10.538	< 2e-16	***
## X31	1.8679	0.3151	5.927	1.93e-08	***
## X32	2.6053	0.3037	8.578	9.35e-15	***
## X33	3.2550	0.2971	10.958	< 2e-16	***
## X34	0.8469	0.2958	2.863	0.004778	**
## X35	2.8325	0.3479	8.143	1.19e-13	***
## X36	1.3446	0.2951	4.557	1.05e-05	***
## X37	3.3069	0.3112	10.627	< 2e-16	***
## X38	2.6245	0.3015	8.704	4.45e-15	***
## X39	2.1381	0.2990	7.152	3.18e-11	***
## X40	1.7628	0.2833	6.223	4.35e-09	***
## X41	2.9588	0.2971	9.960	< 2e-16	***
## X42	2.5143	0.3503	7.177	2.77e-11	***
## X43	3.7129	0.3093	12.003	< 2e-16	***
## X44	1.8393	0.2921	6.296	3.00e-09	***
## X45	1.0386	0.3160	3.287	0.001253	**
## X46	1.9914	0.3203	6.218	4.47e-09	***
## X47	1.3792	0.3116	4.427	1.79e-05	***
## X48	2.0370	0.3089	6.594	6.36e-10	***
## X49	3.0155	0.3122	9.659	< 2e-16	***
## X50	1.6025	0.2964	5.406	2.39e-07	***
## X51	2.3854	0.2870	8.310	4.49e-14	***
## X52	2.2253	0.3012	7.389	8.59e-12	***
## X53	2.2475	0.3058	7.348	1.08e-11	***
## X54	2.0225	0.2994	6.754	2.72e-10	***
## X55	4.3913	0.2938	14.945	< 2e-16	***
## X56	3.5769	0.3291	10.870	< 2e-16	***
## X57	3.9237	0.3292	11.918	< 2e-16	***
## X58	1.6454	0.2899	5.675	6.62e-08	***
## X59	2.3402	0.3190	7.337	1.15e-11	***
## X60	1.4962	0.2991	5.002	1.52e-06	***
## X61	2.0669	0.2987	6.919	1.13e-10	***
## X62	1.8844	0.3186	5.915	2.05e-08	***
## X63	2.9723	0.3184	9.336	< 2e-16	***
## X64	2.0512	0.3040	6.747	2.84e-10	***
## X65	2.9233	0.2974	9.828	< 2e-16	***
## X66	1.9337	0.3042	6.357	2.19e-09	***
## X67	2.9563	0.3009	9.824	< 2e-16	***
## X68	2.3102	0.2997	7.708	1.43e-12	***
## X69	3.4407	0.3172	10.849	< 2e-16	***
## X70	2.1576	0.2982	7.235	2.01e-11	***
## X71	2.9113	0.3272	8.897	1.42e-15	***

```
## X72          4.2092      0.2986  14.096 < 2e-16 ***
## X73          2.8380      0.3109   9.129 3.53e-16 ***
## X74          2.3955      0.2976   8.048 2.05e-13 ***
## X75          2.3246      0.3277   7.093 4.39e-11 ***
## X76          2.1763      0.3306   6.583 6.74e-10 ***
## X77          1.6540      0.3046   5.430 2.14e-07 ***
## X78          3.6930      0.2992  12.342 < 2e-16 ***
## X79          1.7106      0.3304   5.177 6.91e-07 ***
## X80          2.0892      0.2986   6.997 7.39e-11 ***
## X81          1.9015      0.3047   6.240 4.00e-09 ***
## X82          1.7878      0.3225   5.544 1.25e-07 ***
## X83          2.7165      0.3133   8.669 5.46e-15 ***
## X84          2.5323      0.3014   8.401 2.65e-14 ***
## X85          1.2087      0.3019   4.003 9.66e-05 ***
## X86          2.0173      0.3030   6.657 4.58e-10 ***
## X87          2.4892      0.2852   8.728 3.85e-15 ***
## X88          3.5887      0.3180  11.285 < 2e-16 ***
## X89          2.9519      0.2956   9.986 < 2e-16 ***
## X90          1.1933      0.3427   3.482 0.000647 ***
## X91          2.5278      0.2978   8.489 1.58e-14 ***
## X92          1.6659      0.3062   5.441 2.03e-07 ***
## X93          1.7709      0.2868   6.174 5.59e-09 ***
## X94          1.6002      0.2868   5.579 1.06e-07 ***
## X95          1.4424      0.3057   4.719 5.26e-06 ***
## X96          3.0518      0.3230   9.448 < 2e-16 ***
## X97          2.2871      0.3203   7.141 3.37e-11 ***
## X98          2.9399      0.2948   9.971 < 2e-16 ***
## X99          2.2347      0.2986   7.485 5.04e-12 ***
## X100         1.4744      0.3056   4.825 3.32e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 22.08 on 155 degrees of freedom
## Multiple R-squared:  0.9863, Adjusted R-squared:  0.9775
## F-statistic: 111.8 on 100 and 155 DF,  p-value: < 2.2e-16
```

Your comment: We violated the homoscedasticity by making the variance non-constant. Compared to the model trained on the original data, this model has much lower F-statistic and bigger RSE.

Understanding the advantages of shrinkage methods

In this part, we will be modifying the dependent variables X from task 1 by a linear transformation, represented by a square matrix of `n.dimensions` sides. For demonstration, consider that the identity function represented by the identity matrix:

```
M0 <- diag(n.dimensions)      # Makes an identity matrix
X0 <- X %*% M0                 # You can check that X0 == X
noise <- rnorm(n.samples, sd = 8, mean = 0)
Y0 <- (X0 %*% coefs) + noise # We can reuse the noise as it's independent of X.
```

Task 3

In this task, you will show your understanding of the advantages of Ridge by synthesizing data on which they perform the best. You are supposed to do this by linearly transforming the dataset X as in the example above. In other words, you should construct a matrix which if used in place of $M0$ in the above example would make

Ridge perform better than the other two methods. You should not resort to degenerate cases where you would get a warning about using a rank-deficient matrix. Justify your method.

Scoring note: The difference in the RMSE criterion doesn't need to be large or does not need to be present if you are certain it's just an statistical artifact (which you could verify by re-running the tests multiple times or using the LOOCV in `learnAndTest` function). Your design and justification is what matters for the assignment evaluation and the measurements are here only to guide you. We expect this task to be challenging to students, but once you get the right idea, it is possible to implement it with very small amount of code.

```
M1 <- diag(n.dimensions)
# M1 <- diag(x = rnorm(n.dimensions, sd = 0.2), n.dimensions)

# M1.fraction <- 0.1
M1.selection <- sample(n.dimensions, replace = FALSE, size = 10)
M1.idxs <- M1.selection

#1 <- matrix(
#   # runif(n.dimensions^2, min = 0, max = 2),
#   # prob = c(0.7, 0.3)
#   # sample(c(0, 1, 2, 3, 4, 5), replace = TRUE, size = n.dimensions^2),
#   1,
#   nrow = n.dimensions,
#   ncol = n.dimensions,
# )

# M1[row(M1) == col(M1)] <- 0

#
# M1.tmp <- sample(c(0, 1), replace = TRUE, size = n.dimensions * length(M1.idxs), prob = c(0.7, 0.3))
#
# M1[, M1.idxs] <- 0
M1[row(M1) == col(M1)][M1.idxs] <- 100

# M1 <- matrix(
#   # runif(n.dimensions^2, min = 0, max = 2),
#   # sample(c(0, 1), replace = TRUE, size = n.dimensions^2, prob = c(0.7, 0.3)),
#   nrow = n.dimensions,
#   ncol = n.dimensions,
# )
# M1 <- diag(n.dimensions)
# M1.ur <- sample(100, replace = FALSE, size = 50)
# M1.uc <- sample(100, replace = FALSE, size = 50)
# M1[M1.ur, M1.uc] <- sample(coefs, replace = TRUE, size = 50^2)
# M1[col(M1) == row(M1)] <- 1
# M1 <- matrix(
#   1,
#   nrow = n.dimensions,
#   ncol = n.dimensions,
# )
# M1.half <- (n.dimensions / 2)
# M1[, (M1.half + 1):n.dimensions] <- 0
```

```

# M1[1:M1.half, 1:M1.half] <- 0
# M1[col(M1) == row(M1)] <- 0.05
# M1[
#   col(M1) > M1.half &
#   row(M1) > M1.half &
#   col(M1) == row(M1)
# ] <- 1
# M1[1, 1] <- 1

# KEEP THE CODE BELOW
noise <- rnorm(n.samples, sd = 8, mean = 0)
X1 <- X %*% M1
Y1 <- (X1 %*% coefs) + noise

M1.c <- round(cor(X1), 2) > 0.9
test.pca <- prcomp(X1)

summary(test.pca)

```

Ridge

Importance of components:

##	PC1	PC2	PC3	PC4	PC5	PC6	
## Standard deviation	669.975	635.8636	618.1497	612.6362	598.6935	579.4355	
## Proportion of Variance	0.132	0.1189	0.1123	0.1103	0.1054	0.0987	
## Cumulative Proportion	0.132	0.2508	0.3631	0.4735	0.5788	0.6775	
##	PC7	PC8	PC9	PC10	PC11	PC12	
## Standard deviation	555.94339	529.23575	515.19477	489.35120	8.90688	8.69171	
## Proportion of Variance	0.09086	0.08234	0.07803	0.07039	0.00002	0.00002	
## Cumulative Proportion	0.76839	0.85073	0.92875	0.99915	0.99917	0.99919	
##	PC13	PC14	PC15	PC16	PC17	PC18	PC19
## Standard deviation	8.49953	8.41474	8.34254	8.29512	8.13161	7.92884	7.83090
## Proportion of Variance	0.00002	0.00002	0.00002	0.00002	0.00002	0.00002	0.00002
## Cumulative Proportion	0.99922	0.99924	0.99926	0.99928	0.99930	0.99932	0.99933
##	PC20	PC21	PC22	PC23	PC24	PC25	PC26
## Standard deviation	7.81935	7.66864	7.65636	7.58187	7.45721	7.40331	7.32942
## Proportion of Variance	0.00002	0.00002	0.00002	0.00002	0.00002	0.00002	0.00002
## Cumulative Proportion	0.99935	0.99937	0.99939	0.99940	0.99942	0.99944	0.99945
##	PC27	PC28	PC29	PC30	PC31	PC32	PC33
## Standard deviation	7.17902	7.17049	7.07414	7.00771	6.94535	6.89184	6.85866
## Proportion of Variance	0.00002	0.00002	0.00001	0.00001	0.00001	0.00001	0.00001
## Cumulative Proportion	0.99947	0.99948	0.99950	0.99951	0.99952	0.99954	0.99955
##	PC34	PC35	PC36	PC37	PC38	PC39	PC40
## Standard deviation	6.74447	6.67259	6.54965	6.52510	6.45375	6.41294	6.33012
## Proportion of Variance	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001
## Cumulative Proportion	0.99957	0.99958	0.99959	0.99960	0.99962	0.99963	0.99964
##	PC41	PC42	PC43	PC44	PC45	PC46	PC47
## Standard deviation	6.31135	6.16821	6.09280	6.07045	5.99003	5.89599	5.81212
## Proportion of Variance	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001
## Cumulative Proportion	0.99965	0.99966	0.99967	0.99968	0.99970	0.99971	0.99972
##	PC48	PC49	PC50	PC51	PC52	PC53	PC54
## Standard deviation	5.72234	5.70523	5.59733	5.56984	5.43994	5.39251	5.35795
## Proportion of Variance	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001
## Cumulative Proportion	0.99973	0.99973	0.99974	0.99975	0.99976	0.99977	0.99978

	PC55	PC56	PC57	PC58	PC59	PC60	PC61
## Standard deviation	5.26415	5.23008	5.18011	5.11853	5.06531	4.97547	4.91258
## Proportion of Variance	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001
## Cumulative Proportion	0.99979	0.99980	0.99980	0.99981	0.99982	0.99983	0.99983

	PC62	PC63	PC64	PC65	PC66	PC67	PC68
## Standard deviation	4.89528	4.81930	4.78540	4.75813	4.73132	4.62681	4.55507
## Proportion of Variance	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001
## Cumulative Proportion	0.99984	0.99985	0.99985	0.99986	0.99987	0.99987	0.99988

	PC69	PC70	PC71	PC72	PC73	PC74	PC75
## Standard deviation	4.50562	4.43473	4.38530	4.34265	4.26539	4.25630	4.19277
## Proportion of Variance	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001	0.00001
## Cumulative Proportion	0.99988	0.99989	0.99990	0.99990	0.99991	0.99991	0.99992

	PC76	PC77	PC78	PC79	PC80	PC81	PC82	PC83
## Standard deviation	4.13556	4.1026	3.9977	3.9640	3.9326	3.853	3.789	3.733
## Proportion of Variance	0.00001	0.0000	0.0000	0.0000	0.0000	0.000	0.000	0.000
## Cumulative Proportion	0.99992	0.9999	0.9999	0.9999	0.9999	1.000	1.000	1.000

	PC84	PC85	PC86	PC87	PC88	PC89	PC90	PC91	PC92
## Standard deviation	3.634	3.572	3.526	3.38	3.337	3.32	3.217	3.116	3.015
## Proportion of Variance	0.000	0.000	0.000	0.00	0.000	0.00	0.000	0.000	0.000
## Cumulative Proportion	1.000	1.000	1.000	1.00	1.000	1.00	1.000	1.000	1.000

	PC93	PC94	PC95	PC96	PC97	PC98	PC99	PC100
## Standard deviation	2.856	2.821	2.792	2.668	2.626	2.566	2.384	2.216
## Proportion of Variance	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
## Cumulative Proportion	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000

Running these test should now make Ridge perform the best.

```
print(learnAndTest(X1, Y1, "lm"))
```

```
## Linear Regression
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 232, 231, 230, 231, 229, 228, ...
## Resampling results:
##
##   RMSE      Rsquared   MAE
##  10.56176   0.9999935   8.405541
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
print(learnAndTest(X1, Y1, "ridge"))
```

```
## Ridge Regression
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 230, 230, 231, 230, 231, 229, ...
## Resampling results across tuning parameters:
##
```

```
##      lambda  RMSE      Rsquared    MAE
##      0e+00   10.15353  0.9999934    8.182463
##      1e-04   10.19487  0.9999933    8.215248
##      1e-01  460.77793  0.9883638   371.674820
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was lambda = 0.

print(learnAndTest(X1, Y1, "lasso"))

## The lasso
##
## 256 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 231, 232, 231, 230, 231, 228, ...
## Resampling results across tuning parameters:
##
##      fraction  RMSE      Rsquared    MAE
##      0.1       3541.6149  0.4948989  2866.9125
##      0.5       1786.2961  0.9538453  1469.3268
##      0.9       146.2663  0.9990000   118.3576
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was fraction = 0.9.
```

Justification: Ridge works better than OLS and LASSO when there is multicollinearity in the features.

LASSO (OPTIONAL CHALLENGE) This part of the homework is purely optional, but we are eager to see students capable of solving this. Here you should do the same thing as above, but to make perform LASSO the best of the three methods. Although similar in nature, we consider this even more challenging than the above since being unable to modify the underlying coefficients, this may require some deeper considerations to justify the transformation method. It can still be implemented with a few short lines of code, though.

```
# 1. Start with the identity transformation.
M2 <- diag(n.dimensions)
# 2. Let's randomly select 80 % of the features (which we'll make unimportant by scaling the values).
M2.unimportant_fraction <- 0.8
M2.selection <- sample(n.dimensions, replace = FALSE, size = round(M2.unimportant_fraction * n.dimensions))
M2.unimportant_idxes <- M2.selection
# 3. Update the transformation matrix so that it will scale down the selected features when applied.
# Note: We know the magnitude of the coefficients in our case, so we can choose correct scaling factor.
# We can also add the term `(1/coef[M2.unimportant_idxes])` (assuming abs(coef) > 1)
# to further normalize the scaling factor.
M2[, M2.unimportant_idxes] <- M2[, M2.unimportant_idxes] *
  0.001 *
  (1 / coef[M2.unimportant_idxes])
# M2[,M2.unimportant_idxes] <- M2[,M2.unimportant_idxes] * 0.001 # also works great

# KEEP THE CODE BELOW
noise <- rnorm(n.samples, sd = 8, mean = 0)
X2 <- X %*% M2
```

```
Y2 <- (X2 %*% coefs) + noise
```

Running these test should now make LASSO perform the best.

```
t31 <- learnAndTest(X2, Y2, "lm")
t32 <- learnAndTest(X2, Y2, "ridge")
t33 <- learnAndTest(X2, Y2, "lasso")
t32_c <- as.matrix(predict(t32$finalModel, s = t32$bestTune[1, "fraction"], type = "coef", mode = "fraction"))
t33_c <- as.matrix(predict(t33$finalModel, s = t33$bestTune[1, "fraction"], type = "coef", mode = "fraction"))
```

Justification: Generally, LASSO performs better when the response is a function of only a relatively small number of predictors. We can adjust the data by a linear transformation to make it the case. We create a scaling matrix that scales down a selected subset of features (making them unimportant) which also affects the linear regression coefficients. This in turn pushes LASSO to perform variable selection (forcing some of the coefficient estimates to be exactly equal to zero when lambda is sufficiently large). We ran the tests multiple times for our transformation and in all instances, the LASSO was the best (lowest RMSE, with a difference of around 1).